

Testing In The Cloud

How we reduced testing time by 50%
and eliminated waiting time entirely



Kenneth Falck <kennu@sc5.io>

4.11.2014

Presentation Overview

Case Veikkaus.fi

- The Problem
- Background
- Our Solution
- How Did It Go?

CASE VEIKKAUS.FI



Lotto Viking Lotto Eurojackpot Jokeri Naapurit Keno eBingo Arvat Syke Pore Porukkapelit
Urheilu Pitkäveto Tulosveto Moniveto Voittajavedot Live-veto Vakio Moniveikkaus

Tulokset Asiakaspalvelu
Ostoskori Asiakasedut

**Lotto 1 300 000 €**

Loton potti tuplattuna 2 600 000 € Jokeri 1 000 000 €

Viikko 44 - peliaikaa su 22.00 asti

Siirry pelikuponkiin Ohjeet ja säännöt Tulokset

Pikapeli

Lotto

- 2 +

Tuplaus +0,25 €/rivi

Ei Kyllä

UUTTA

1: 4 12 17 19 26 35 37 2: 2 4 9 11 17 18 33

Jokeri

- 1 +

Kesto

1 arvonta

1: 5 2 0 7 0 2 0

Maksa 4,00 €

Porukkapelit

Järjestelmä 8

8 riviä, 1/5 osuus, 1 rivi Lauantai-Jokeria

3 4 11 17 24 26 27 28

Maksa 2,40 €

Harava 18

600 riviä, 1/10 osuus, 2 riviä Lauantai-Jokeria

2 3 6 9 11 14 16 17 18 19 20 24 28 29
30 33 34 37

Maksa 75,40 €

Näytä kaikki

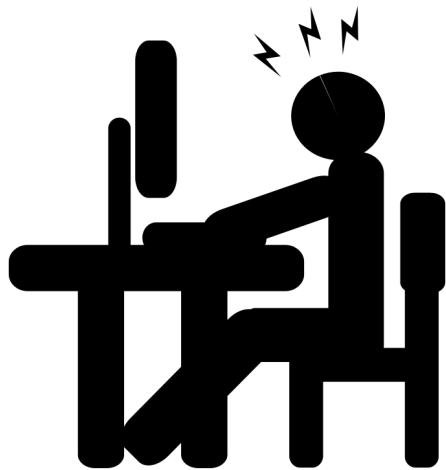
VEIKKAUS.FI

Front End Web Project

- HTML/JavaScript based online lottery games
- Large (10+) development team
- Everybody working simultaneously on the same web project
- Every change must be thoroughly tested

THE PROBLEM

Testing Is Slow



- Developers need to run tests several times a day
- One test run can take 30 minutes
 - First wait 10 min for a free slot
 - Then wait 20 min for tests to run
- As a result, developers become less productive and time is wasted

BACKGROUND

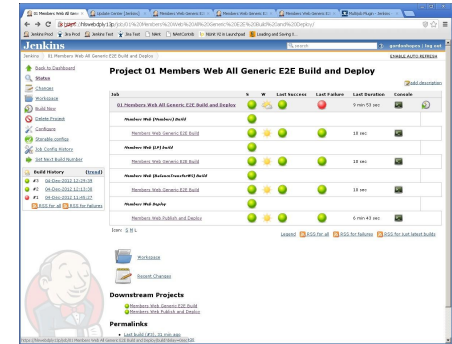
What is Jenkins?



Jenkins is a testing and deployment automation platform.

Developers use Jenkins to test a new version of their project, and to deploy it to production.

Developers don't need to know any details - They click “Test” and wait for the green light. 🟢



Jenkins Web UI

What are Branches?



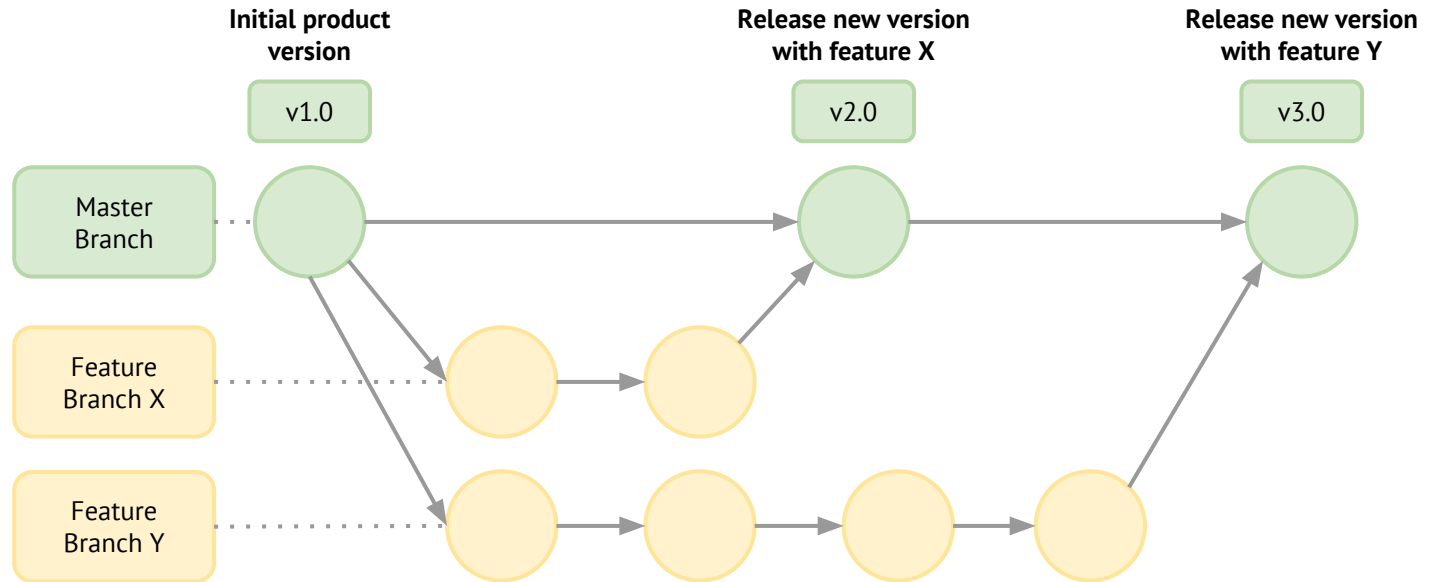
Master is what gets published to production



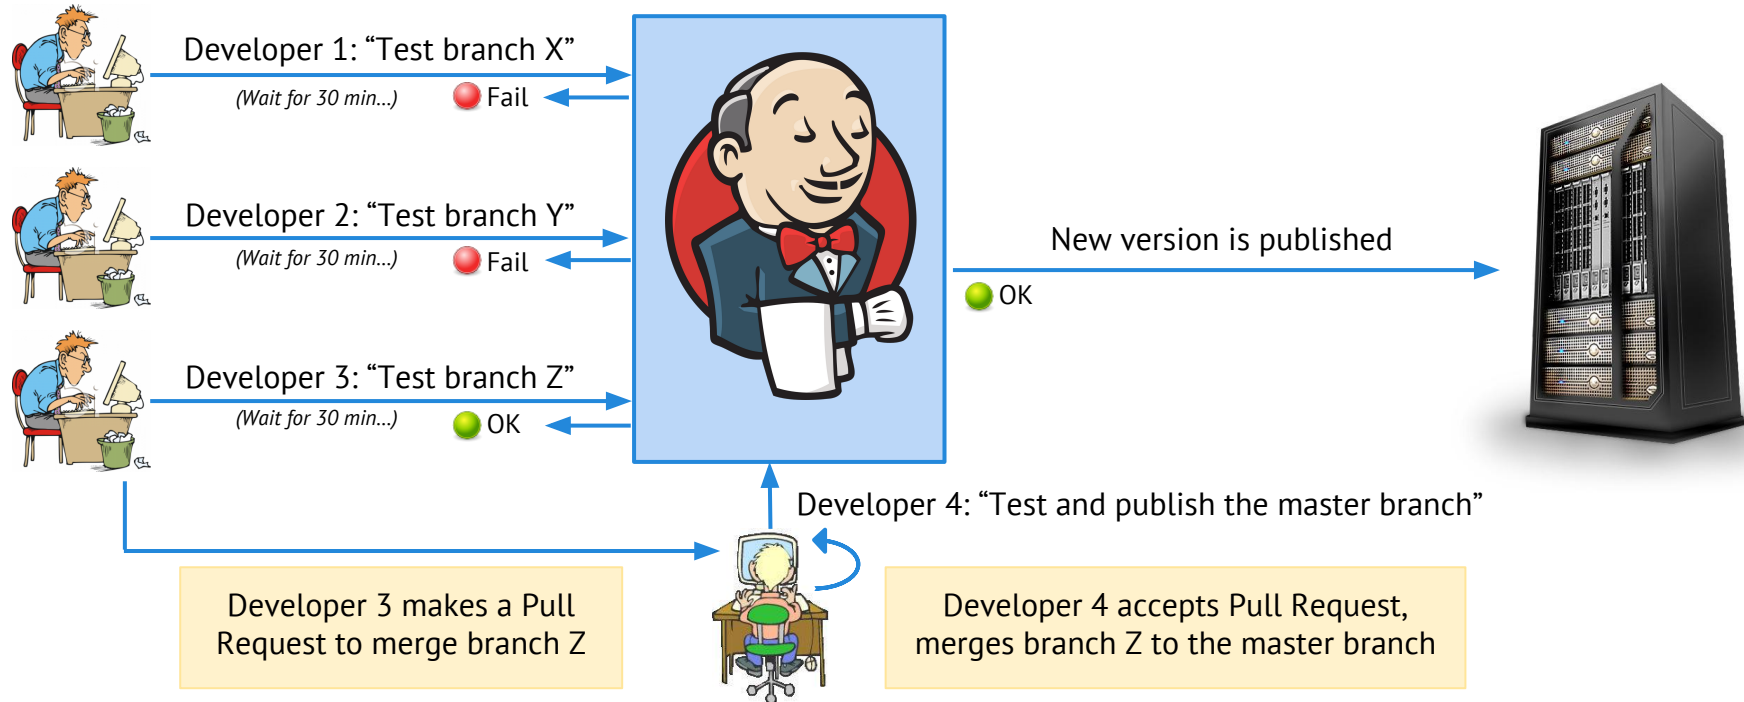
Developer 1 works on feature X in her own branch



Developer 2 works on feature Y in her own branch



How Development Works



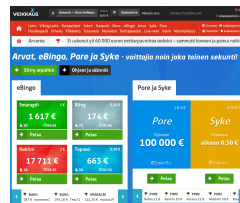
Why Is Testing Slow?



PhantomJS is the virtual web browser that runs our automatic JavaScript tests without human interaction.

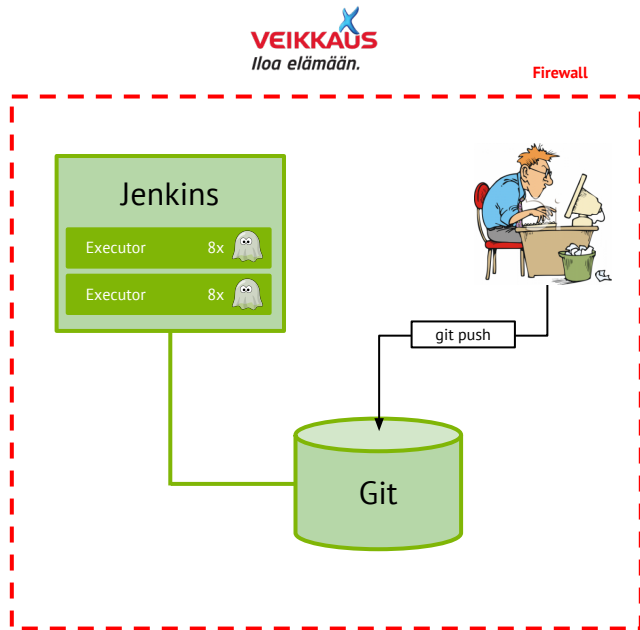
One PhantomJS requires one CPU core. To run 8 tests in parallel, 8 cores are needed. This makes testing 8 times faster.

To maximize testing speed, 400+ cores would be needed (every test would run simultaneously).



**OUR SOLUTION
TO SPEED THINGS UP**

Original Setup at Veikkaus



A physical 16-core Jenkins server (no cloud)

- Scarce resource shared by many coders, no scaling
- Only two parallel job executors, 8 PhantomJS each
- Full test suite needs 400+ PhantomJS runs

Long waiting and running times

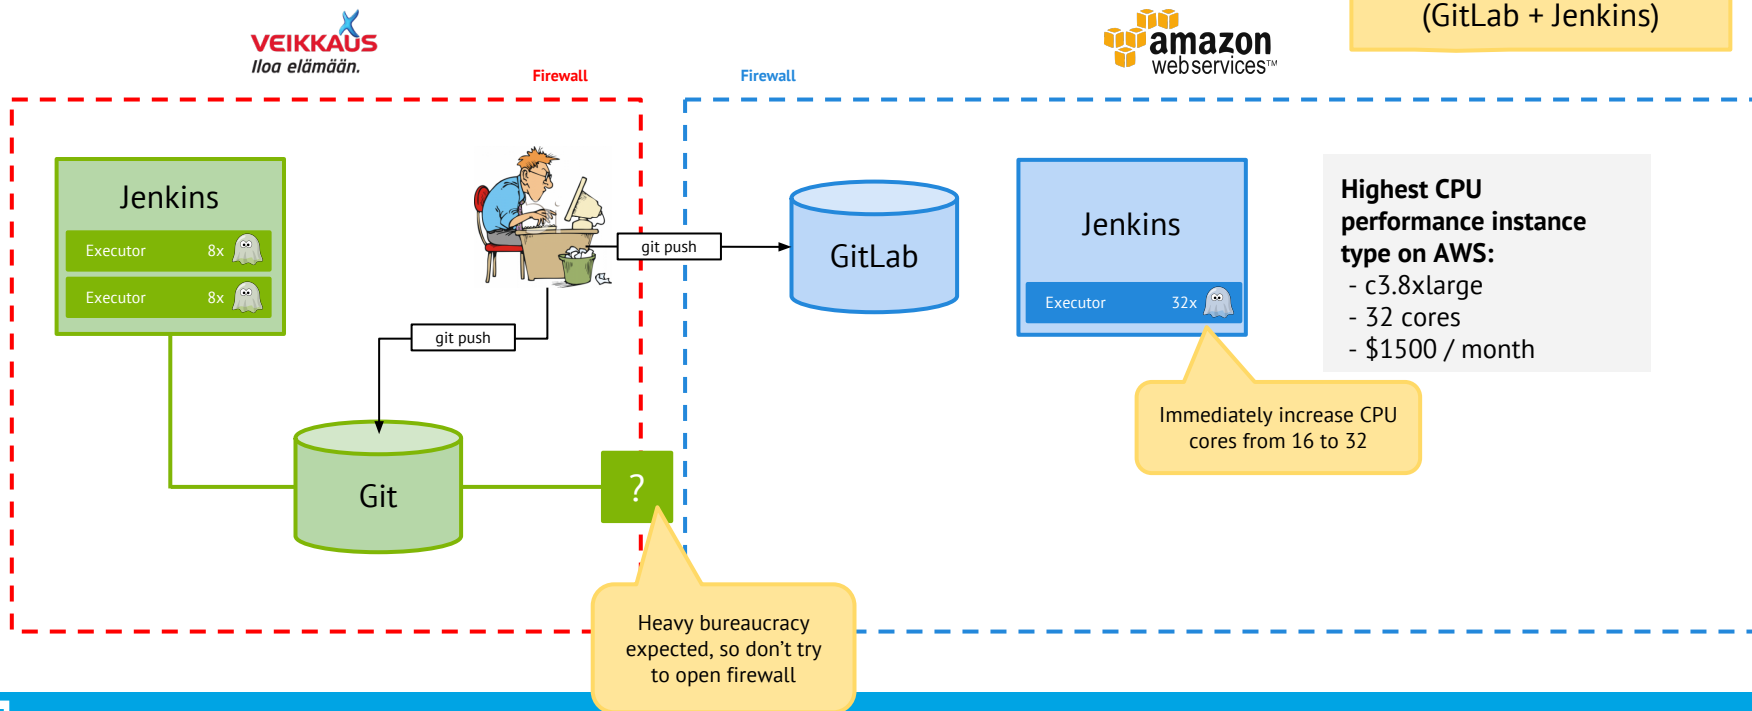
- First wait 10 minutes for the previous job to end
- Then wait 20 minutes for your own job to finish

How could we improve this?

- Let's move it to the cloud (Amazon AWS)
- We made a proposal of a scalable architecture and an analysis of the security implications (accepted)

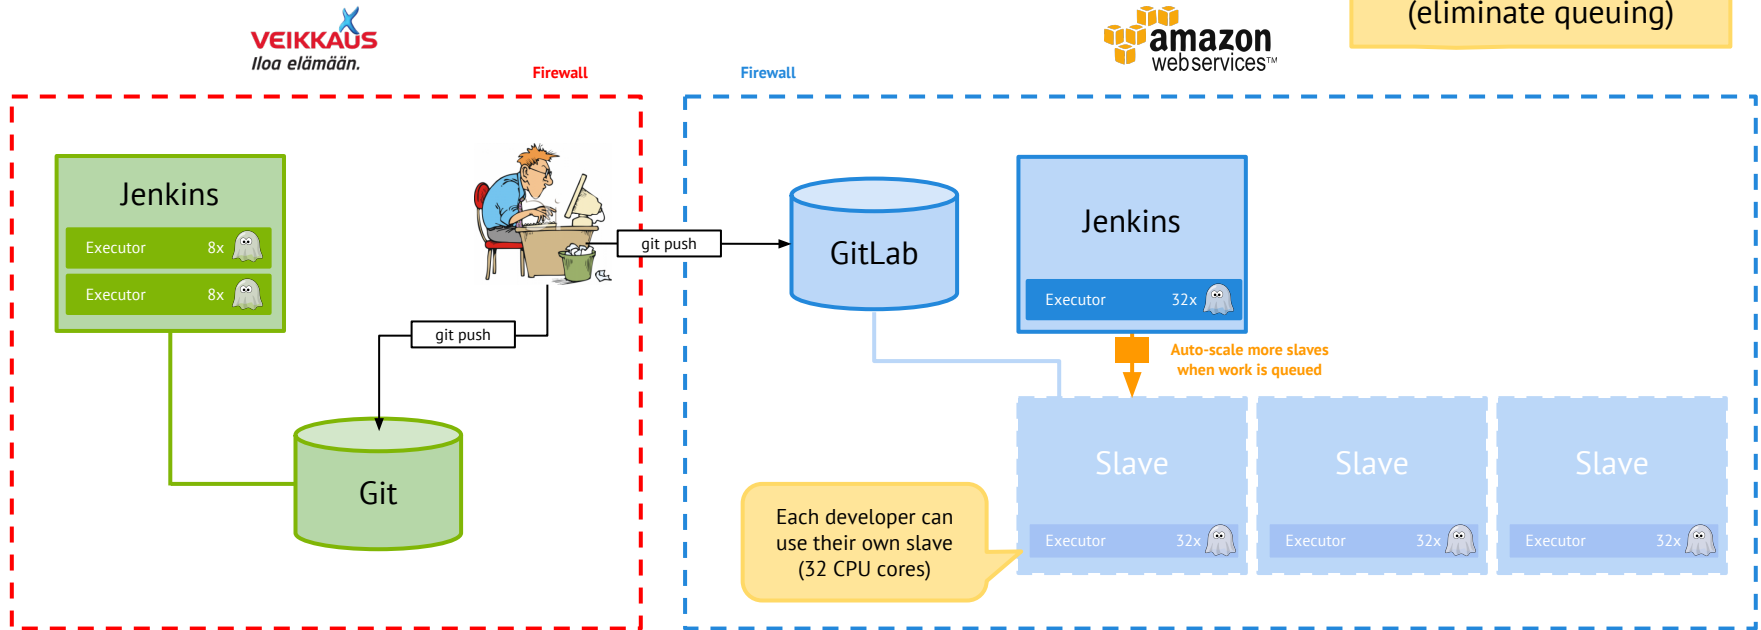
Phase 1: Jenkins on Amazon

Goal: Setup basic cloud infrastructure (GitLab + Jenkins)



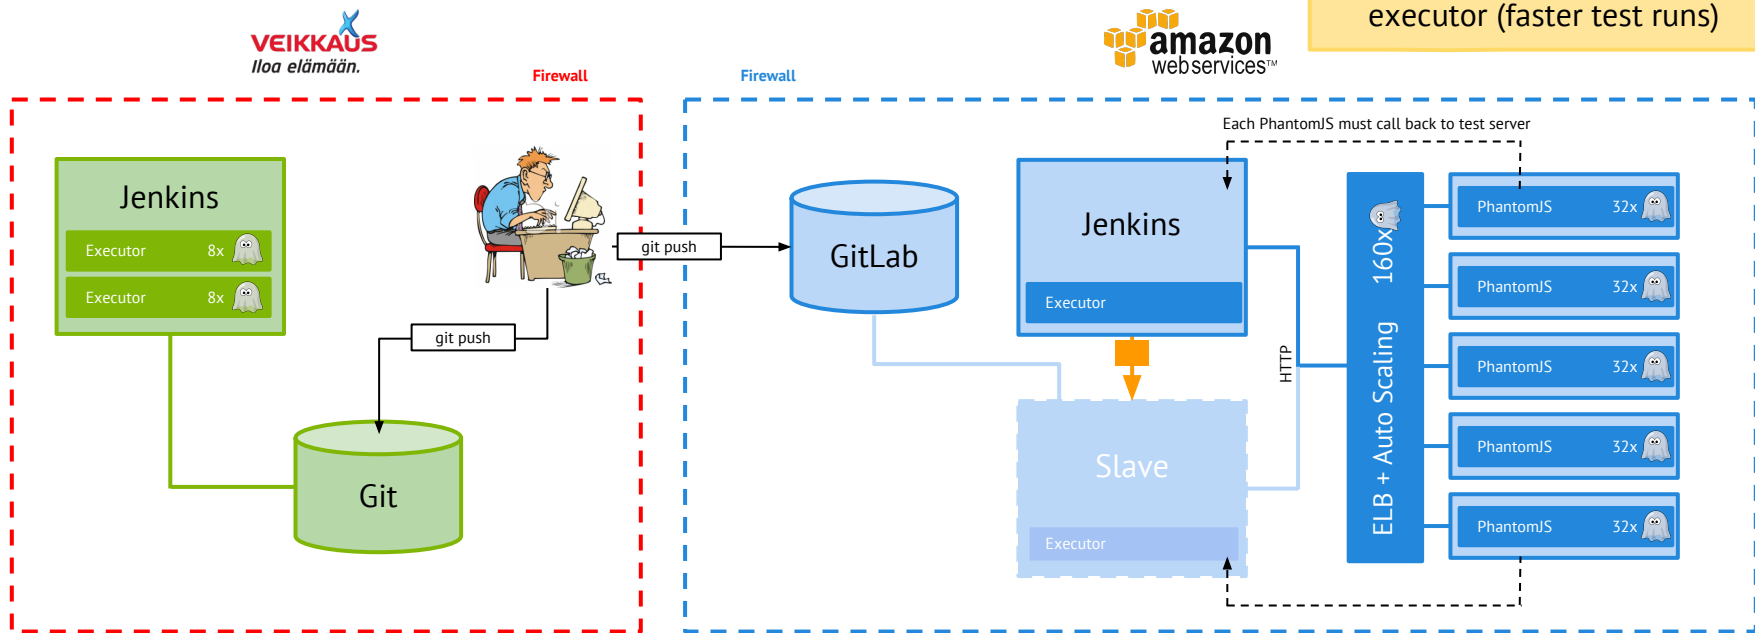
Phase 2: Scalable Jenkins

Goal: Scale out to N parallel job executors (eliminate queuing)



Phase 3: PhantomJS Cluster

Goal: Scale out to more than 32 parallel Phantoms per job executor (faster test runs)



HOW DID IT GO?

Jenkins & GitLab on AWS

MISSION
ACCOMPLISHED

Continue using
internal Jenkins
for deployments

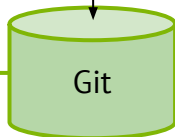
VEIKKAUS
Iloa elämään.

Firewall

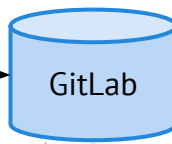
Firewall

amazon
webservices™

Developers use AWS
Jenkins for testing their
own branches



git push



GitLab has built-in
user management



Jenkins has built-in
user management

Auto-scale more slaves
when work is queued

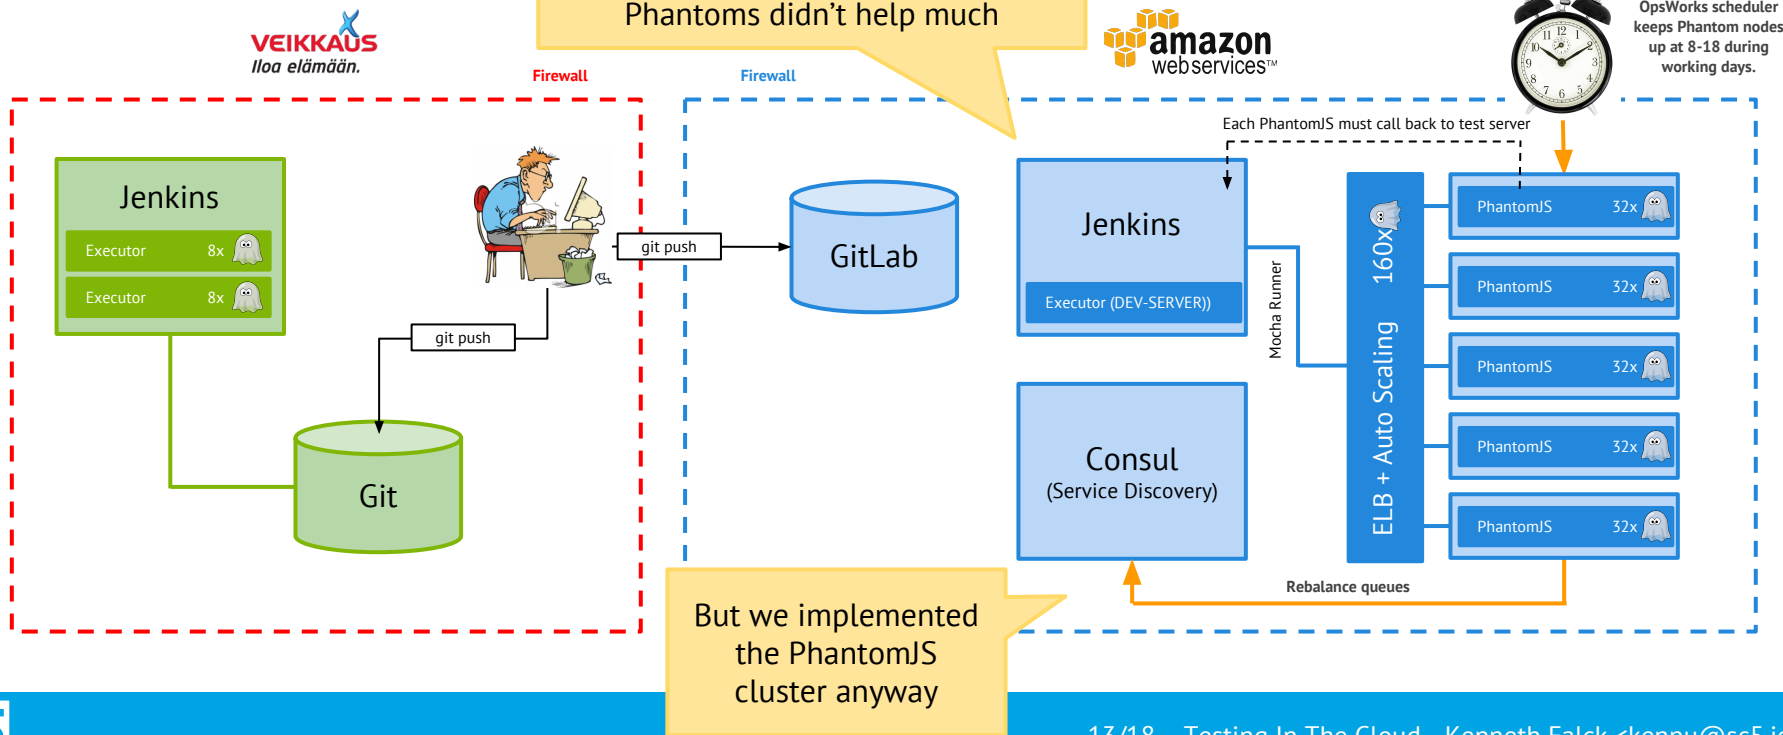


"Caffeine" job keeps slave up 8-18 every working day, starts more slaves if all are busy (terminate after 30 min idle)


PhantomJS Cluster



We found that this part (dev-server) didn't scale well, so adding more Phantoms didn't help much



Clean AWS Architecture



OpsWorks

STACK 1

Chef Recipe

Jenkins

Run Command Stack Settings Delete Stack

A stack represents a collection of EC2 instances and related AWS resources that have a common purpose and that you want to manage collectively. Within a stack, you use layers to define the configuration of your instances and use apps to specify the code you want to deploy. [Learn more.](#)




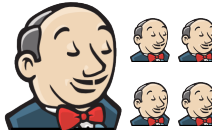
Instances

Layers 1

Jenkins

Apps 0

No apps. [Add an app.](#)



Jenkins + slaves GitLab Consul Manager UI



OpsWorks

STACK 2

Chef Recipe

Phantom Cluster

Run Command Stack Settings Delete Stack

A stack represents a collection of EC2 instances and related AWS resources that have a common purpose and that you want to manage collectively. Within a stack, you use layers to define the configuration of your instances and use apps to specify the code you want to deploy. [Learn more.](#)

Instances

Layers 1

Node.js App Server

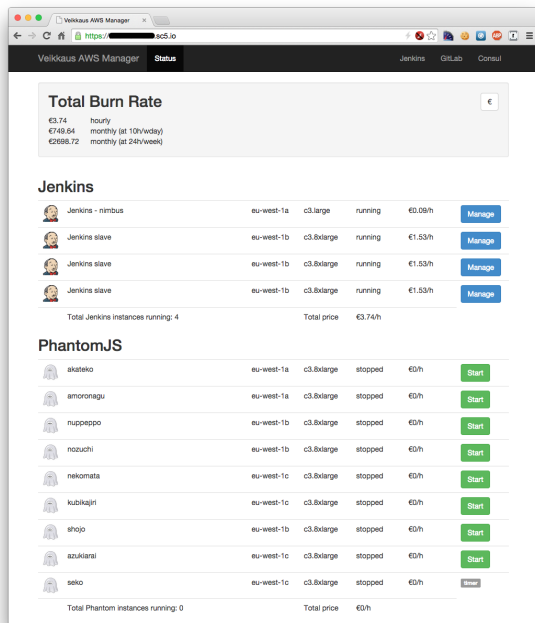
Apps 1

Phantom Server deploy



PhantomJS Cluster Nodes (custom Node.js app)

AWS Management UI



Custom Angular.js/Node.js App

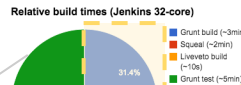
- Simple UI to check cloud status
- Calculate monthly cost of running instances
- Manually start/stop PhantomJS cluster nodes

Test Speeds and AWS Pricing

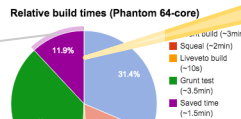
Setup	Total time	Test time	At 10h/workday	Always on 24/7
Jenkins 32-core	~10 min	5 min 7 s	\$469 / month	\$1,464 / month
Phantom 32-core	~10 min	4 min 38 s	\$852 / month	\$2,840 / month
Phantom 64-core	9 min 12 s	3 min 35 s	\$1,234 / month	\$4,217 / month
Phantom 96-core	9 min 33 s	3 min 55 s	\$1,617 / month	\$5,594 / month
Phantom 128-core	9 min 17 s	3 min 38 s	\$2,007 / month	\$6,770 / month
Phantom 160-core	9 min 17 s	3 min 38 s	\$2,381 / month	\$8,347 / month

Not exact measurements, because PhantomJS often crashes and causes variations in test time

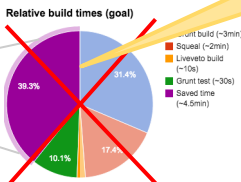
But still, end results were **50%** of the original time (~20 min)



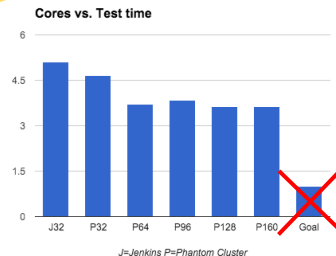
Half of the spent time was fixed (build and init time)



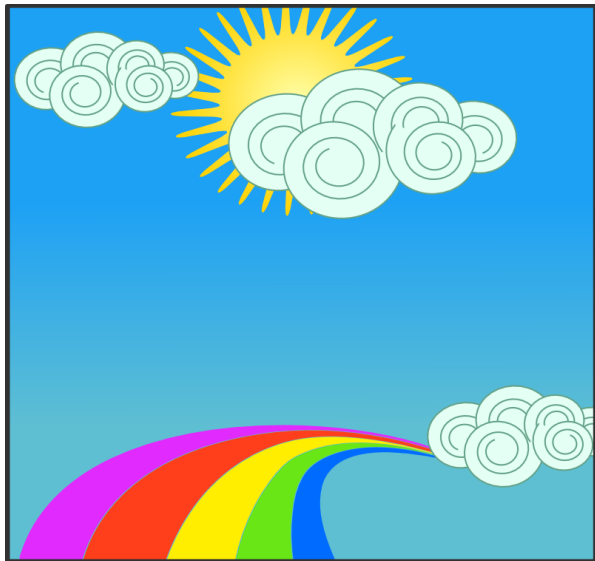
Violet sector = saved time



Goal couldn't be fully reached because dev-server did not scale



Bonus: Tests Are More Robust



Fixed many issues in tests

- Catch failures and error situations
- Eliminate dependencies on timing

Fixed PhantomJS crashing

- Retry tests when PhantomJS crashes

Made tests more universal

- Run locally on developer computers
- Run remotely in the cloud

Future: How To Be Even Faster



Get 192+ cores to work

- Tests started to fail after adding 160 cores
- Bottleneck in the development server
- The first test immediately generates 5000+ HTTP requests

Split test files that are too big

- If a test file takes 2 minutes to run, that sets the hard limit on speed (parallelization doesn't help any further)

Optimize PhantomJS execution

- Reduce startup time (pool / prefork Phantom processes)
- Disable image, font and CSS requests

THANKS!

Twitter: @kennu

Email: kennu@sc5.io